

DSP Kaggle Project II

Titanic: Machine Learning from Disaster

Group 2

Brian Lai
Kirill Evseev
Alejandra Bermudez

California State University, Northridge

NSF support for the DSP with Career Support and Connections to Industry Program
Award # 1842386

June 20, 2019

Hello Data

The screenshot displays a multitasking environment. The primary window is Microsoft Excel, titled 'train - Excel', which contains a dataset of Titanic passengers. The data is organized into columns: PassengerId, Survived, Pclass, Name, Sex, Age, SibSp, Parch, Ticket, Fare, Cabin, and Embarked. The spreadsheet lists 32 passengers, with details such as their survival status, social class, and family size.

Overlaid on the right side of the Excel window is a web browser displaying a Python script. The visible code includes:

```
tree.py in _validate_X_predic  
  
validation.py in check_array  
es, ensure_min_features, warn  
  
rane.py in _setitem__(self,  
    ...)
```

The bottom of the screen shows the Windows taskbar with several open applications, including 'Cleaning and Pr...ipynb', 'Cleaning the Dat...ipynb', 'ds19-master (3)ip', 'ds19-master (2)ip', 'Project 2(ka)...ipynb', and 'ds19-master (1)ip'. The system clock in the bottom right corner indicates the time is 8:27 PM on 6/20/2019.

A look at the Data

In [20]:


```
# Mean: survival rate  
# Count: total observations  
# Sum: people survived  
# Sex: 0 female, 1 male  
data[['Sex', 'Survived']].groupby(['Sex'], as_index=False).agg(['mean', 'count', 'sum'])
```

Out[20]:

	Survived		
	mean	count	sum
Sex			
0	0.754789	261	197
1	0.205298	453	93

Another Look at the Data

```
# survival rate per age
group_by_age = pd.cut(train["Age"], np.arange(0, 90, 10))
age_grouping = train.groupby(group_by_age).mean()
age_grouping['Survived'].plot.bar()
plt.ylabel('Survival Rate')
plt.grid(c="lightblue")
```



How we Cleaned the Data

First Attempt

```
In [12]: 1 # dropping the data that does not contribute
          2 data = train.drop(columns=['Cabin','Embarked','Name','Ticket','Fare'],axis=0)
```

```
[Feature]
Cabin      687
Age        177
Embarked    2
Fare        0
Ticket     0
Parch      0
SibSp      0
Sex         0
Name        0
Pclass     0
Survived    0
PassengerId 0
dtype: int64
```

Cleaning the Data

```
1 # Cleaning the data: first attempt
```

```
In [70]: 1 train3=pd.DataFrame(train)
```

```
In [71]: 1 train3.drop(['Name', 'SibSp', 'Parch', 'Ticket', 'Fare', 'Cabin', 'Embarked'], axis=1, inplace=True)
```

```
In [72]: 1 train3=train3.dropna(axis=0)
```

```
In [73]: 1 train3['Age']=np.where(train3['Age'].between(0,10),0,train3['Age'])
2 train3['Age']=np.where(train3['Age'].between(10.5,30),1,train3['Age'])
3 train3['Age']=np.where(train3['Age'].between(30.5,60),2,train3['Age'])
4 train3['Age']=np.where(train3['Age'].between(60.5,100),3,train3['Age'])
```

```
In [75]: 1 train3=pd.get_dummies(train3,columns=['Sex'])
2 train3=pd.get_dummies(train3,columns=['Age'])
3 train3=pd.get_dummies(train3,columns=['Pclass'])
```

Cleaned Data

In [88]: 1 train3[:10]

Out[88]:

	PassengerId	Survived	Sex_female	Sex_male	Age_0.0	Age_1.0	Age_2.0	Age_3.0	Pclass_1	Pclass_2	Pclass_3
0	1	0	0	1	0	1	0	0	0	0	1
1	2	1	1	0	0	0	1	0	1	0	0
2	3	1	1	0	0	1	0	0	0	0	1
3	4	1	1	0	0	0	1	0	1	0	0
4	5	0	0	1	0	0	1	0	0	0	1
6	7	0	0	1	0	0	1	0	1	0	0
7	8	0	0	1	1	0	0	0	0	0	1
8	9	1	1	0	0	1	0	0	0	0	1
9	10	1	1	0	0	1	0	0	0	1	0
10	11	1	1	0	1	0	0	0	0	0	1

Second Cleaning Attempt (SCA)

```
In [30]: 1 tr4=pd.DataFrame(train) #going with tr4 instead of train4
```

```
In [31]: 1 tr4.drop(['Ticket', 'Fare', 'Cabin', 'Embarked'], axis=1, inplace=True)
```

```
In [32]: 1 data=tr4
2 titles = {"Mr": 1, "Miss": 2, "Mrs": 3, "Master": 4, "Rare": 5}
3 for dataset in data:
4     dataset['Title'] = dataset.Name.str.extract('([A-Za-z]+)\.', expand=False)
5     dataset['Title'] = dataset['Title'].replace(['Lady', 'Countess', 'Capt', 'Col', 'Don', 'Dr', \
6         'Major', 'Rev', 'Sir', 'Jonkheer', 'Dona'], 'Rare')
7     dataset['Title'] = dataset['Title'].replace('Mlle', 'Miss')
8     dataset['Title'] = dataset['Title'].replace('Ms', 'Miss')
9     dataset['Title'] = dataset['Title'].replace('Mme', 'Mrs')
10    dataset['Title'] = dataset['Title'].map(titles)
11    dataset['Title'] = dataset['Title'].fillna(0)
12 tr4 = tr4.drop(['Name'], axis=1)
13 #taken from https://towardsdatascience.com/predicting-the-survival-of-titanic-passengers-30870ccc7e8
```

```
In [33]: 1 titleNull=tr4[tr4['Title']==0] #there is nobody with a null title
2 len(titleNull)
```

```
Out[33]: 0
```

Trying to find out the mean age of those who are possibly children or adults with the thinking spouse/sib<1 would mean siblings and 0<parents/children<=2 would probably mean parents if the first was satisfied

```
In [34]: 1 tr4pc=tr4[tr4['Parch']<=2] #tr4pc is possible child
2 tr4pc=tr4pc[tr4pc['Parch']>0] #trying no parents when # is in front
3 tr4pc=tr4pc[tr4pc['SibSp']>1]
4 #last might mess things up since the assumption is that spouse with no kids but it could be someone with only one sibling
5 len(tr4pc)
```

```
Out[34]: 55
```

```
In [35]: 1 tr4pa=pd.concat([tr4,tr4pc,tr4pc]).drop_duplicates(keep=False)#tr4pa is possible adult
2 len(tr4pa)
```


Getting Some Results (GSR)

```
3 meanAgeMr=mr["Age"].mean()
4 medAgeMr=mr["Age"].median()
5 print(meanAgeMr,medAgeMr,len(mr))
6
7 miss=tr4LA[tr4LA['Title']==2]
8 #miss=miss[miss['Parch']==0] #likely older
9 meanAgeMs=miss["Age"].mean()
10 medAgeMs=miss["Age"].median()
11 print(meanAgeMs,medAgeMs,len(miss))
12
13 mrs=tr4LA[tr4LA['Title']==3]
14 meanAgeMrs=mrs["Age"].mean()
15 medAgeMrs=mrs["Age"].median()
16 print(meanAgeMrs,medAgeMrs,len(mrs))
17
18 rare=tr4LA[tr4LA['Title']==5]
19 meanAgeRare=rare["Age"].mean()
20 medAgeRare=rare["Age"].median()
21 print(meanAgeRare,medAgeRare,len(rare))
```

```
32.56106870229008 30.0 509
27.68617021276596 25.5 121
35.898148148148145 35.0 125
45.54545454545455 48.5 23
```

Making an Informed Guess (MIG)


In [45]:

```
1 #Replace age with mean of each larger group rounded to the near  
2 tr4LCh=tr4LCh.fillna(9.5)  
3 tr4LA=tr4LA.fillna(33)
```

Resulting Age Distribution (RAD)

```
# Survival rate per age of combined  
group_by_age = pd.cut(train_combine["Age"], np.arange(0, 90, 8))  
age_grouping = train_combine.groupby(group_by_age).mean()  
age_grouping['Survived'].plot.bar()  
plt.grid(c="lightblue")  
plt.ylabel('Survival Rate')
```

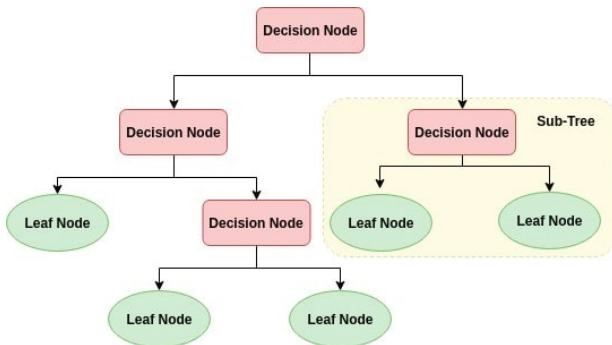
Text(0, 0.5, 'Survival Rate')



What is our Model?

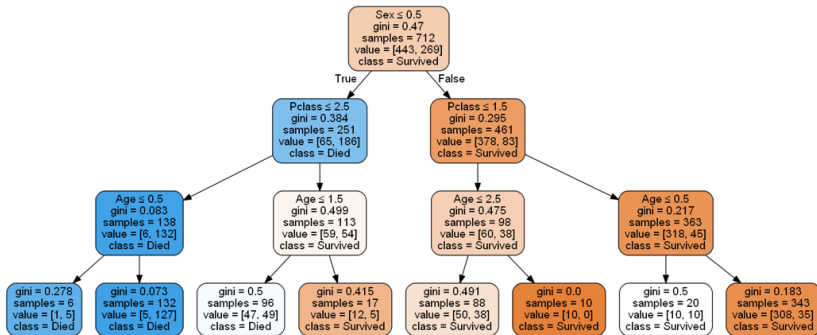
What exactly is a decision tree?

- A Decision Tree is a support tool that uses a flowchart-like model of decisions and potential consequences
- A supervised learning algorithm that works with both continuous and categorical variables
- It is a tool that only contains conditional control statements
- Each leaf node is a class label which is the outcome



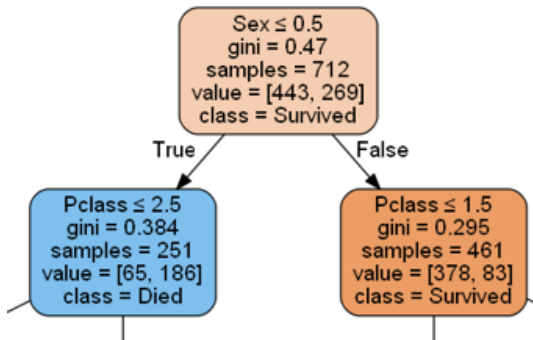
Model: Decision Tree

- Each branch is the outcome of the test
- It's rarely ever balanced
- A depth of 3:
Sex, Pclass, and Age



Model: Decision Tree

- Most important factor is Sex, it splits off to determine survival/death outcome
- Blue boxes = most likely to die
- Orange boxes = most likely to survive



Fitting the Model

```
In [78]: 1 from sklearn import tree
        2 from sklearn.model_selection import train_test_split
        3 from sklearn.metrics import accuracy_score
```

```
In [79]: 1 X=train3.values
        2 Y=train3['Survived'].values
```

```
In [80]: 1 X=np.delete(X,1,axis=1)
```

```
In [81]: 1 X=np.delete(X,0,axis=1)
```

```
In [85]: 1 XTRAIN, XTEST, YTRAIN, YTEST=train_test_split(X,Y)
        2 nsplit=1000
        3 depth=1
```


```
In [87]: 1 while (depth < 7):
        2     errs=[]
        3     for j in range(nsplit):
        4         XTRAIN, XTEST, YTRAIN, YTEST=train_test_split(X,Y)
        5         DT=tree.DecisionTreeClassifier(max_depth=depth)
        6         DT.fit(XTRAIN,YTRAIN)
        7         YP=DT.predict(XTEST)
        8         errs.append(1-accuracy_score(YTEST,YP))
        9     print("Decision Tree Depth = %d mean error = %7.6f SD=%7.6f\"
        10         %(depth,np.mean(errs),np.std(errs)))
        11     depth = depth + 1
```

```
Decision Tree Depth = 1 mean error = 0.220464 SD=0.026738
Decision Tree Depth = 2 mean error = 0.213799 SD=0.026424
Decision Tree Depth = 3 mean error = 0.209737 SD=0.026049
Decision Tree Depth = 4 mean error = 0.204296 SD=0.025150
Decision Tree Depth = 5 mean error = 0.209514 SD=0.024734
Decision Tree Depth = 6 mean error = 0.209944 SD=0.024799
```

Determining the Optimal Depth

```
Decision Tree Depth = 1 mean error = 0.220028 SD=0.026047  
Decision Tree Depth = 2 mean error = 0.212240 SD=0.027499  
Decision Tree Depth = 3 mean error = 0.207553 SD=0.026333  
Decision Tree Depth = 4 mean error = 0.205508 SD=0.024900  
Decision Tree Depth = 5 mean error = 0.208112 SD=0.025645  
Decision Tree Depth = 6 mean error = 0.208754 SD=0.025076
```

```
In [99]: 1 plt.plot((range(1,7)), errors, color='red')  
2 plt.xlabel("DEPTH")  
3 plt.ylabel("MEAN ERROR")  
4 plt.grid()  
5 plt.gcf().set_size_inches(10,6)
```



Results: Initial Attempt

From our first attempt:
Score = 0.55980

Results: Second Attempt

From our second attempt:

Score = 0.78468

An increase of 0.22488!

References

<https://towardsdatascience.com/predicting-the-survival-of-titanic-passengers-30870ccc7e8?gi=339b274bc177>

<https://www.datacamp.com/community/tutorials/decision-tree-classification-python>

<https://blog.patricktriest.com/titanic-machine-learning-in-python/>

